

freiheit.com
technologies
hamburg / lisbon

FOUNDED 1999
OVER 120 FULL STACK ENGINEERS
SINCE 2018 NEW OFFICE IN LISBON

```

R9  0x7ffff7f9ace0 -> 0x7ffff799e78 -> 0x7ffff7ea7440 (.__cxixabiv1::__class_type_info::~__class_type_info()) <- endbr64
R10 0x7ffff7ab4f40 <- stosq qword ptr [rdi], rax /* 0x7ffff7ab4f40 */
R11 0x246
R12 0x401090 (._start) <- endbr64
R13 0x7fffffffef30 <- 0x1
R14 0x0
R15 0x0
RBP 0x7fffffffef950 <- 0x4141414141414141 ('AAAAAAAA')
RSP 0x7fffffffef930 <- 0x4141414141414141 ('AAAAAAAA')
RIP 0x4011fb (main+74) <- lea rsi, [rip + 0xe2e]

```

[DISASM]

```

0x4011e3 <main+50>    call    0x401070

0x4011e8 <main+55>    lea     rax, [rbp - 0x20]
0x4011ec <main+59>    mov     rsi, rax
0x4011ef <main+62>    lea     rdi, [rip + 0x2f8a] <0x404180>
0x4011f6 <main+69>    call    0x401030

▶ 0x4011fb <main+74>    lea     rsi, [rip + 0xe2e]
0x401202 <main+81>    lea     rdi, [rip + 0x2e57] <0x404060>
0x401209 <main+88>    call    0x401030

0x40120e <main+93>    mov     rcx, rcx
0x401211 <main+96>    lea     rax, [rbp - 0x20]
0x401215 <main+100>   mov     rsi, rax

```

Hacking like in the 90s

Binary Exploitation Workshop

[STACK]

```

00:0000 | rsp 0x7fffffffef930 <- 0x4141414141414141 ('AAAAAAAA')
... ↓

```

[BACKTRACE]

```

▶ f 0      4011fb main+74

```

\$(whoami)

- Vitali Henne (30)
- Software Engineer at freiheit.com
- CTF player with
 - Cyclopropenylidene
 - SauerCloud
 - KITCTF in the past
- I like to break things :>



WTF is a CTF?

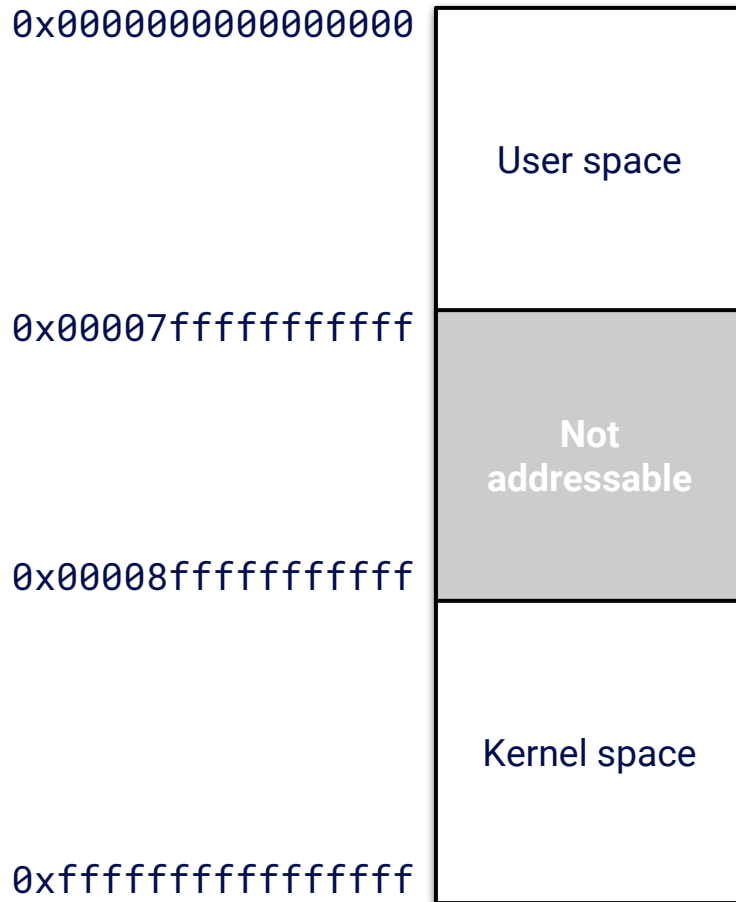
- **Practical infosec competition**
 - Usually organized by teams participating in such competitions for other teams
- **Different categories:**
 - Web
 - Reverse engineering
 - Binary exploitation
 - Cryptography
 - ...
- **Duration varies between 12h - 72h**
- **Goal is to retrieve a “flag” - a piece of information hidden in the challenge:**
 - E.g. In binary exploitation the flag is somewhere on the server and you want to exploit the binary to leak this data

WTF is a CTF?

- **Practical infosec competition**
 - Usually organized by teams participating in such competitions for other teams
- **Different categories:**
 - Web
 - Reverse engineering
 - **Binary exploitation**
 - Cryptography
 - ...
- **Duration varies between 12h - 72h**
- **Goal is to retrieve a “flag” - a piece of information hidden in the challenge:**
 - E.g. In binary exploitation the flag is somewhere on the server and you want to exploit the binary to leak this data

Memory Layout

- Each process is assigned a distinct virtual address space by the kernel.
- Divided into multiple regions with different access permission:
 - Readable
 - Writeable
 - Executable
- Example: x86_64 address space on Linux



Mem. Layout - Text Segment

0x0000000000000000

- Also called code segment
- Contains machine code of the binary
- Readable | Executable

```
pwndbg> x/8i main
0x4011c1 <main>:      push    rbp
0x4011c2 <main+1>:    mov     rbp, rsp
0x4011c5 <main+4>:    sub     rsp, 0x110
0x4011cc <main+11>:  mov     DWORD PTR [rbp-0x4], 0x2a
```

0x00007fffffff

Text

Mem. Layout - Data Segment

0x0000000000000000

- Initialized data of global and static variables
- Fixed size, known on compile time
- Readable
- Readable | Writeable

```
#include <iostream>

char buf[] = "Hello World!";
int main(int argc, char* argv[]) {
    std::cout << buf << std::endl;
}
```

0x00007fffffff

Text

Data

Mem. Layout - BSS Segment

0x0000000000000000

- Uninitialized global and static variables
- Whole segment initialized with 0 by kernel
- Fixed size, known on compile time
- **Readable** | **Writeable**

```
#include <iostream>
#include <string.h>

char buf[16];
int main(int argc, char* argv[]) {
    strcpy(buf, "Hello World!");
    std::cout << buf << std::endl;
}
```

0x00007fffffff

Text

Data

BSS

Mem. Layout - Stack

- Located in “high memory” area
- Grows towards lower memory addresses
- Local variables of each function
- Readable | Writeable | ~~Executable~~

```
#include <iostream>

int main(int argc, char* argv[]) {
    char buf[] = "Hello World!";
    std::cout << buf << std::endl;
}
```

0x0000000000000000

Text

Data

BSS

Stack


0x00007fffffffffff

Mem. Layout - Heap

- Located behind the BSS segment
- Grows towards higher memory addresses
- Managed by a heap allocator:
 - Dynamically allocates chunks of memory
 - Chunk size may be defined on runtime
- **Readable** | **Writable**

```
// ...  
int main(int argc, char* argv[]) {  
    char* buf = malloc(16);  
    strcpy(buf, "Hello World!");  
    std::cout << buf << std::endl;  
    free(buf);  
}
```

0x0000000000000000



0x00007fffffffffff


Mem. Layout - Shared Libs

0x0000000000000000

- Memory mapping segment
- Readable
- Readable | Writeable
- Readable | Executable

```
pwndbg> vmmap
LEGEND: STACK | HEAP | CODE | DATA | RWX | RODATA
0x555555554000 0x555555555000 r--p 1000 0 /home/vagrant/my_binary
0x555555555000 0x555555556000 r-xp 1000 1000 /home/vagrant/my_binary
0x555555556000 0x555555557000 r--p 1000 2000 /home/vagrant/my_binary
0x555555557000 0x555555558000 r--p 1000 2000 /home/vagrant/my_binary
0x555555558000 0x555555559000 rw-p 1000 3000 /home/vagrant/my_binary
0x7ffff7de4000 0x7ffff7e06000 r--p 22000 0 /usr/lib/libc-2.28.so
0x7ffff7e06000 0x7ffff7f51000 r-xp 14b000 22000 /usr/lib/libc-2.28.so
0x7ffff7f51000 0x7ffff7f9d000 r--p 4c000 16d000 /usr/lib/libc-2.28.so
0x7ffff7f9d000 0x7ffff7f9e000 --p 1000 1b9000 /usr/lib/libc-2.28.so
0x7ffff7f9e000 0x7ffff7fa2000 r--p 4000 1b9000 /usr/lib/libc-2.28.so
0x7ffff7fa2000 0x7ffff7fa4000 rw-p 2000 1bd000 /usr/lib/libc-2.28.so
0x7ffff7fa4000 0x7ffff7faa000 rw-p 6000 0
0x7ffff7fce000 0x7ffff7fd1000 r--p 3000 0 [vvar]
0x7ffff7fd1000 0x7ffff7fd3000 r-xp 2000 0 [vdso]
0x7ffff7fd3000 0x7ffff7fd5000 r--p 2000 0 /usr/lib/ld-2.28.so
0x7ffff7fd5000 0x7ffff7ff4000 r-xp 1f000 2000 /usr/lib/ld-2.28.so
0x7ffff7ff4000 0x7ffff7ffc000 r--p 8000 21000 /usr/lib/ld-2.28.so
0x7ffff7ffc000 0x7ffff7ffd000 r--p 1000 28000 /usr/lib/ld-2.28.so
0x7ffff7ffd000 0x7ffff7ffe000 rw-p 1000 29000 /usr/lib/ld-2.28.so
0x7ffff7ffe000 0x7ffff7fff000 rw-p 1000 0
0x7ffff7fff000 0x7ffff7fff000 rw-p 21000 0 [stack]
```

0x00007ffffffffff



The Stack Frame

- Holds local variables
- Allocate a frame for each function called
 - Required size is known on compile time because all local variables are known

```
void do_stuff() {  
    char buf[] = "Hello World!";  
    int x = 1337;  
    void* ptr = malloc(23);  
    float y = 23.23;  
}
```

0x00007fffffffe8f0

23.23

0x5390040137d

1337

Hello World!

???

???

0x00007fffffffe920

Previous Stack
Frame

The Call Stack

- **ip**: points to the next to-be executed instruction
- **sp**: points to the current end of the stack

```
void f2() {  
    printf("Called f2\n");  
    return;  
}
```

```
void f1() {  
    printf("Called f1\n");  
    f2();  
    return;  
}
```

```
ip → int main() {  
    f1();  
    printf("done\n");  
}
```

Main() stack frame is allocated upon invocation




The Call Stack

- **ip**: points to the next to-be executed instruction
- **sp**: points to the current end of the stack

```
void f2() {  
    printf("Called f2\n");  
    return;  
}  
  
void f1() {  
    printf("Called f1\n");  
    f2();  
    return;  
}  
  
int main() {  
    f1();  
    printf("done\n");  
}
```

ip →

**f1() stack frame is allocated
when it's called by main()**




The Call Stack

- **ip**: points to the next to-be executed instruction
- **sp**: points to the current end of the stack

```
void f2() {  
    printf("Called f2\n");  
    return;  
}  
  
void f1() {  
    printf("Called f1\n");  
    f2();  
    return;  
}  
  
int main() {  
    f1();  
    printf("done\n");  
}
```

ip →

**f2() stack frame is allocated
when it's called by f1()**




The Call Stack

- **ip**: points to the next to-be executed instruction
- **sp**: points to the current end of the stack

```
void f2() {  
    printf("Called f2\n");  
    return;  
}  
  
void f1() {  
    printf("Called f1\n");  
    f2();  
    return;  
}  
  
int main() {  
    f1();  
    printf("done\n");  
}
```

ip →

What happens when f2() returns?



The Call Stack

- **ip**: points to the next to-be executed instruction
- **sp**: points to the current end of the stack


```
void f2() {  
    printf("Called f2\n");  
    return;  
}  
  
void f1() {  
    printf("Called f1\n");  
    f2();  
    return;  
}  
  
int main() {  
    f1();  
    printf("done\n");  
}
```

Diagram illustrating the state of the program:

- The instruction pointer (**ip**) is shown pointing to the first instruction of `f2()` (the `printf` statement).
- A dashed arrow indicates the previous instruction pointer position, pointing to the `f2();` call inside `f1()`.
- A red arrow points to the `f2();` call inside `f1()`, indicating the instruction that was just executed.

What happens when `f2()` returns?

=> Continue execution of `f1()`
=> Restore `f1()` stackframe



The Call Stack

Restoration of previous state

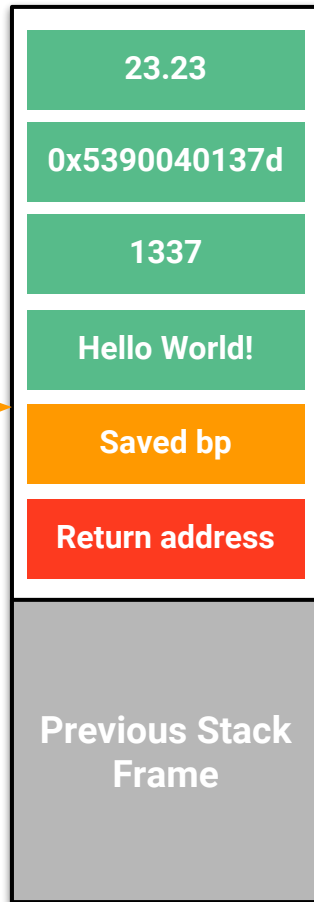
- a) Return address
- b) Previous stack pointer
 - i) Use a “stack base pointer” **bp** pointing to the start of the stack frame

```
void do_stuff() {  
    char buf[] = "HeLo World!";  
    int x = 1337;  
    void* ptr = malloc(23);  
    float y = 23.23;  
}
```

sp → 0x00007fffffffe8f0

bp →

0x00007fffffffe920




Function Prologue

```
(1) call foo (push rip, jmp foo)
(2) push rbp
(3) mov rbp, rsp
(4) sub rsp 32
```

0x00007fffffffef8f0

sp → 0x00007fffffffef920
bp →



Function Prologue

- The call instruction pushes **ip** on the stack

→ (1) `call foo` (push `rip`, jmp `foo`)

(2) `push rbp`

(3) `mov rbp, rsp`

(4) `sub rsp 32`

0x00007fffffffef0

sp →

Return address

sp →

0x00007fffffffef20

1337

bp →

Saved bp

Return address

Function Prologue

- The call instruction pushes **ip** on the stack
- Save the previous base pointer

```
→ (1) call foo (push rip)
  (2) push rbp
  (3) mov rbp, rsp
  (4) sub rsp 32
```

0x00007fffffffef0

sp →

Saved bp

Return address

sp →

0x00007fffffffef20

bp →

1337

Saved bp

Return address

Function Prologue

- The call instruction pushes **ip** on the stack
- Save the previous base pointer
- Set new base pointer

(1) `call foo` (push `rip`, jmp `foo`)

(2) `push rbp`

→ (3) `mov rbp, rsp`

(4) `sub rsp 32`

0x00007fffffffef80

bp → sp →

Saved bp

Return address

sp →

0x00007fffffffef20

bp →

1337

Saved bp

Return address

Function Prologue

- The call instruction pushes **ip** on the stack
- Save the previous base pointer
- Set new base pointer
- Make room on the stack for local args

(1) `call foo` (push `rip`, jmp `foo`)

(2) `push rbp`

(3) `mov rbp, rsp`

(4) `sub rsp 32`

0x00007fffffffef0

sp →

bp →

Saved bp

Return address

sp →

0x00007fffffffef20

bp →

1337

Saved bp

Return address

Function Epilogue

```
(1) mov rsp, rbp  
(2) pop rbp  
(3) ret (pop rip)
```

0x00007fffffffef0

sp →

bp →

Return address

sp →

0x00007fffffffef20

bp →

1337

Saved bp

Return address

...

Function Epilogue

- “Cleanup” local arguments

```
→ (1) mov rsp, rbp  
  (2) pop rbp  
  (3) ret (pop rip)
```

0x00007fffffffef0

sp → bp →

Saved bp

Return address

sp →
0x00007fffffffef20

bp →

1337

Saved bp

Return address

Function Epilogue

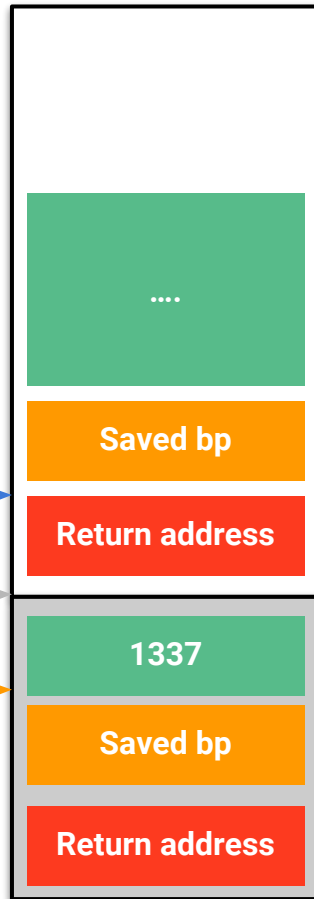
- “Cleanup” local arguments
- Restore previous base pointer

```
→ (1) mov rsp, rbp  
  (2) pop rbp  
  (3) ret (pop rip)
```

0x00007fffffffef0

sp →
0x00007fffffffef20

bp →



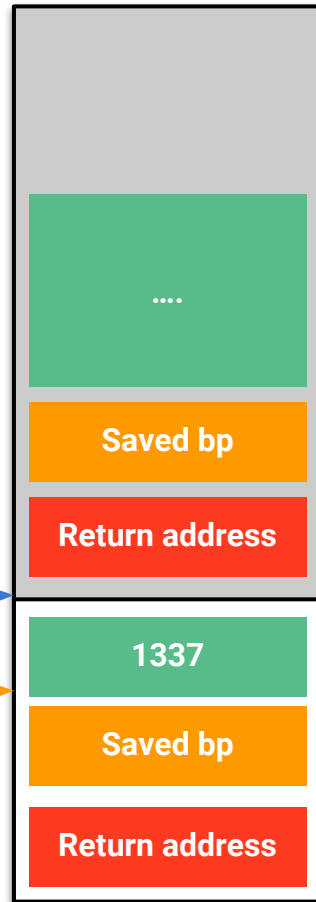
Function Epilogue

- “Cleanup” local arguments
- Restore previous base pointer
- The ret instruction pops **ip** from the stack

```
→ (1) mov rsp, rbp  
  (2) pop rbp  
  (3) ret (pop rip)
```


0x00007fffffffef80

sp → 0x00007fffffffef920
bp →



x86 Assembly 101 - Registers

- Lower parts of general purpose registers can be accessed individually
- x64 extends registers to 64 bit



Current issue : #49 | Release date : 1996-11-08 | Editor : daemon9

Author : Aleph1

.oO Phrack 49 Oo.


Volume Seven, Issue Forty-Nine

File 14 of 16

BugTraq, r00t, and Underground.Org
bring you

XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
Smashing The Stack For Fun And Profit
XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX

by Aleph One
aleph1@underground.org



0x00007fffffffe8f0

Exploiting Stack Overflows


Goal:

Change control flow to achieve arbitrary code execution

Execution will continue at the address stored here!



0x00007fffffffe920




Exploiting Stack Overflows

buf →

What happens if we were to overflow a buffer on the stack?

```
void fn(char *input) {  
    char buf[16];  
    strcpy(buf, input);  
}
```



Exploiting Stack Overflows

What happens if we were to overflow a buffer on the stack? “strcpy” has no check for length!


```
void fn(char *input) {  
    char buf[16];  
    strcpy(buf, input);  
}
```

```
fn("AAAAAAAAAAAAAAAAAAAAAAAAA...")
```


buf →



1. **Get control flow**
2. **Redirect to your code**
3. **???**
4. **Profit**

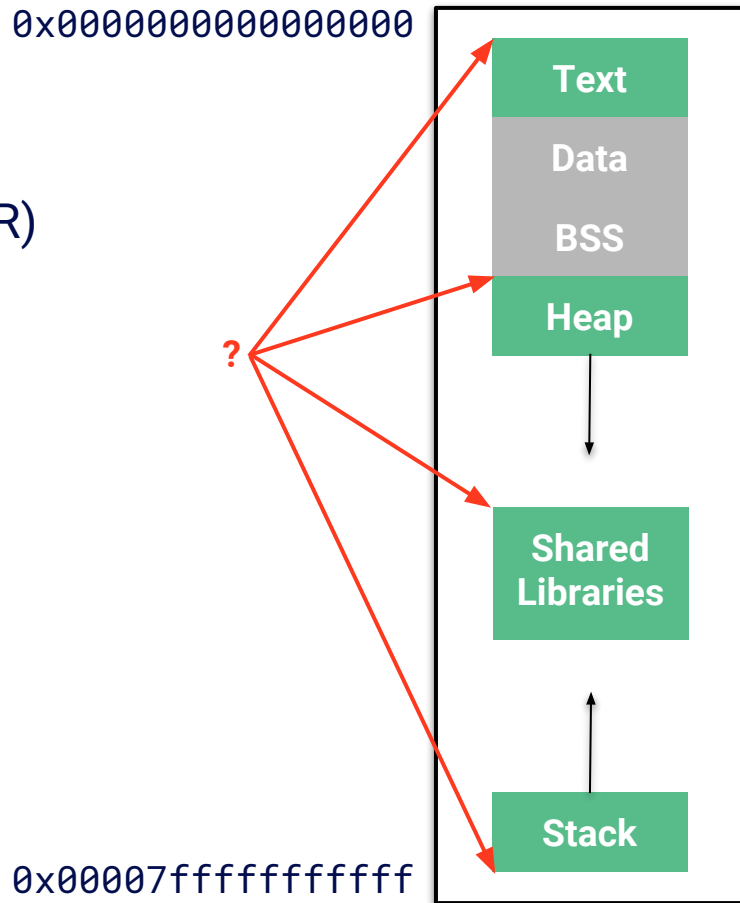


Demo time!!!



Mitigations

- Address Space Layout Randomization (ASLR)
 - Randomize locations of different memory regions
 - => Need an information leak to bypass




Mitigations

- Address Space Layout Randomization (ASLR)
 - Randomize locations of different memory regions
 - => Need an information leak to bypass
- Non-Executable Stack (NX)
 - Readable | Writeable | ~~Executable~~
 - => Use return-oriented programming (ROP)

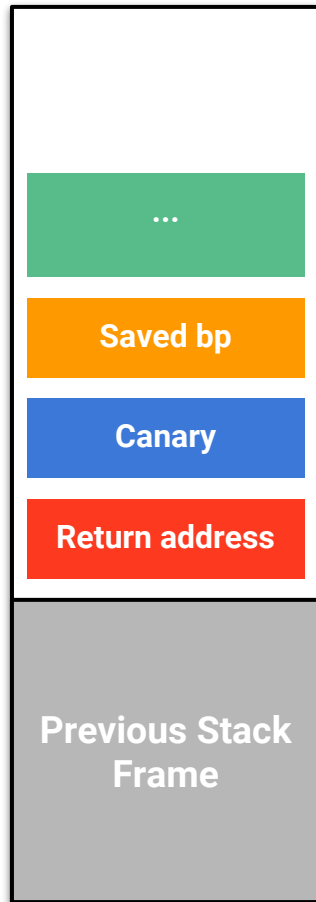
0x0000000000000000

0x00007fffffffffff




Mitigations

- Address Space Layout Randomization (ASLR)
 - Randomize locations of different memory regions
 - => Need an information leak to bypass
- Non-Executable Stack (NX)
 - Readable | Writeable | ~~Executable~~
 - => Use return-oriented programming (ROP)
- Stack Canaries
 - Place a random value right before the return address
 - Compare to its initial value before returning
 - Terminate program if they don't match
 - => Either need to leak the canary or have non-linear overflow




Return-Oriented Programming

- If the nx bit is set, we can't execute code we write :(



Return-Oriented Programming


- If the nx bit is set, we can't execute code we write :(
- => Reuse existing code



Return-Oriented Programming

- If the nx bit is set, we can't execute code we write :'(
- => Reuse existing code
- To call `system("/bin/sh")` on x64
 - Have a pointer to `"/bin/sh"` in `RDI`
 - Call `system` function
- Reminder: Function epilogue:

```
(1) mov rsp, rbp
(2) pop rbp
(3) ret (pop rip)
```




Return-Oriented Programming

Goal:

- Have a pointer to “/bin/sh” in RDI
- Call system function

```
(1) mov rsp, rbp  
(2) pop rbp  
(3) ret (pop rip)
```



Return-Oriented Programming

Goal:

- Have a pointer to “/bin/sh” in RDI
- Call system function

=> Find a code block with **pop rdi; ret** (assume its address is 0x1337)

```
(1) mov rsp, rbp
(2) pop rbp
(3) ret (pop rip)
```

AAAAAAAAAA
AAAAAAAAAA
AAAAAAAAAA
AAAAAAAAAA
AAAAAAAAAA
AAAAAAAAAA
AAAAAAAAAA

0x1337

Return-Oriented Programming

Goal:

- Have a pointer to “/bin/sh” in RDI
- Call system function

=> Find a code block with **pop rdi; ret** (assume its address is 0x1337)

```
(1) mov rsp, rbp
(2) pop rbp
(3) ret (pop rip)
```

sp →

AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA

0x1337

Return-Oriented Programming

Goal:

- Have a pointer to “/bin/sh” in RDI
- Call system function

=> Find a code block with **pop rdi; ret** (assume its address is 0x1337)

=> Write address of “/bin/sh” after it (assume its address is 0x4242)

```
(1) mov rsp, rbp
(2) pop rbp
(3) ret (pop rip)
```

sp →

AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA

0x1337

0x4242

Return-Oriented Programming

Goal:

- Have a pointer to “/bin/sh” in RDI
- Call system function

=> Find a code block with **pop rdi; ret** (assume its address is 0x1337)

=> Write address of “/bin/sh” after it (assume its address is 0x4242)

```
(1) mov rsp, rbp
(2) pop rbp
(3) ret (pop rip)
```

sp →

AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA

0x1337

0x4242

Return-Oriented Programming

Goal:

- Have a pointer to “/bin/sh” in RDI
- Call system function

=> Find a code block with **pop rdi; ret** (assume its address is 0x1337)

=> Write address of “/bin/sh” after it (assume its address is 0x4242)

=> Write address of system after it

```
(1) mov rsp, rbp
(2) pop rbp
(3) ret (pop rip)
```

sp →

AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA
AAAAAAAA

0x1337

0x4242

system


@LateNightSeth



BOOM. EASY AS THAT.

Resources

- <https://ctftime.org> <= All you need to know about upcoming CTFs
- <http://overthewire.org> <= Lots of wargames
- <https://picoc.tf> <= Very beginner friendly CTF, running all year long



HACKERS WANTED.

jobs@freiheit.com